

Stacey Patrick Ken Ben - Customer Churn Version 5

December 6, 2018

1 Customer Churn Dataset

Group 4 - Patrick Codrington, Stacey Jovicic, Ben Polasek, Kenneth Brandt The initial work surrounded loading the data, analyzing the data and preparing the attributes that would be used for modeling.

2 PART 1 - Data Preparation and Pre-prediction Analysis

2.1 A) Initial data load

```
In [1]: #Import Required Libraries

#numpy and pandas for data investigation and cleanup
import numpy as np
import pandas as pd

#pprint
import pprint

#matplotlib and seaborn for visuals
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
plt.style.use('seaborn-whitegrid')

#sklearn for machine learning
from sklearn import tree
#from sklearn.linear_model import LogisticRegression
#from sklearn.svm import SVC, LinearSVC
#from sklearn.ensemble import RandomForestClassifier
#from sklearn.neighbors import KNeighborsClassifier
#from sklearn.naive_bayes import GaussianNB
#from sklearn.linear_model import Perceptron
#from sklearn.linear_model import SGDClassifier
#from sklearn.tree import DecisionTreeClassifier
```

```
In [2]: #Import the customer_churn.csv dataset
df = pd.read_csv('customer_churn.csv')

#Verify import was successful and check for missing values
#Confirmed and there is no missing data
df.info()

#See column names
#print(df.columns.values)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
State                3333 non-null object
Account Length      3333 non-null int64
Area Code           3333 non-null int64
Phone               3333 non-null object
Int'l Plan          3333 non-null object
VMail Plan          3333 non-null object
VMail Message       3333 non-null int64
Day Mins            3333 non-null float64
Day Calls           3333 non-null int64
Day Charge          3333 non-null float64
Eve Mins            3333 non-null float64
Eve Calls           3333 non-null int64
Eve Charge          3333 non-null float64
Night Mins          3333 non-null float64
Night Calls         3333 non-null int64
Night Charge        3333 non-null float64
Intl Mins           3333 non-null float64
Intl Calls          3333 non-null int64
Intl Charge         3333 non-null float64
CustServ Calls      3333 non-null int64
Churn?              3333 non-null object
dtypes: float64(8), int64(8), object(5)
memory usage: 546.9+ KB
```

2.2 B) What are the attribute type? (e.g. categorical, ordinal or quantitative)

```
In [3]: #Preview the data, to determine types
df.head()
```

```
Out[3]:
```

	State	Account Length	Area Code	Phone	Int'l Plan	VMail Plan	\
0	KS	128	415	382-4657	no	yes	
1	OH	107	415	371-7191	no	yes	
2	NJ	137	415	358-1921	no	no	
3	OH	84	408	375-9999	yes	no	

4	OK	75	415	330-6626	yes	no
---	----	----	-----	----------	-----	----

	VMail Message	Day Mins	Day Calls	Day Charge	...	Eve Calls	\
0	25	265.1	110	45.07	...	99	
1	26	161.6	123	27.47	...	103	
2	0	243.4	114	41.38	...	110	
3	0	299.4	71	50.90	...	88	
4	0	166.7	113	28.34	...	122	

	Eve Charge	Night Mins	Night Calls	Night Charge	Intl Mins	Intl Calls	\
0	16.78	244.7	91	11.01	10.0	3	
1	16.62	254.4	103	11.45	13.7	3	
2	10.30	162.6	104	7.32	12.2	5	
3	5.26	196.9	89	8.86	6.6	7	
4	12.61	186.9	121	8.41	10.1	3	

	Intl Charge	CustServ Calls	Churn?
0	2.70	1	False.
1	3.70	1	False.
2	3.29	0	False.
3	1.78	2	False.
4	2.73	3	False.

[5 rows x 21 columns]

What are the categorical attributes?

Categorical: State, Area Code, Phone, Int'l Plan (binary), VMail Message (binary), Churn (class variable).

What are the numerical attributes?

Continuous: Account Length, Day Mins, Day Charge, Eve Mins, Eve Charge, Night Mins, Night Charge, Intl Mins, Intl Charge Discrete: VMail Message, Day Calls, Eve Calls, Night Calls, Intl Calls, CustServ Calls

3 C) Find the max, min, mean and standard deviation of each attribute

Use describe to first gather this data for the categorical values and convert area code to object as this field is not numeric.

```
In [4]: df['Area Code'] = df['Area Code'].astype(object)
df.describe(include=['O'])
```

```
Out[4]:
```

	State	Area Code	Phone	Int'l Plan	VMail Plan	Churn?
count	3333	3333	3333	3333	3333	3333
unique	51	3	3333	2	2	2
top	WV	415	381-2745	no	no	False.
freq	106	1655	1	3010	2411	2850

Use describe again to this time gather the data for numerical values.

```
In [5]: df.describe()
```

```
Out [5]:
```

	Account Length	VMail Message	Day Mins	Day Calls	Day Charge \
count	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000
mean	101.064806	8.099010	179.775098	100.435644	30.562307
std	39.822106	13.688365	54.467389	20.069084	9.259435
min	1.000000	0.000000	0.000000	0.000000	0.000000
25%	74.000000	0.000000	143.700000	87.000000	24.430000
50%	101.000000	0.000000	179.400000	101.000000	30.500000
75%	127.000000	20.000000	216.400000	114.000000	36.790000
max	243.000000	51.000000	350.800000	165.000000	59.640000

	Eve Mins	Eve Calls	Eve Charge	Night Mins	Night Calls \
count	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000
mean	200.980348	100.114311	17.083540	200.872037	100.107711
std	50.713844	19.922625	4.310668	50.573847	19.568609
min	0.000000	0.000000	0.000000	23.200000	33.000000
25%	166.600000	87.000000	14.160000	167.000000	87.000000
50%	201.400000	100.000000	17.120000	201.200000	100.000000
75%	235.300000	114.000000	20.000000	235.300000	113.000000
max	363.700000	170.000000	30.910000	395.000000	175.000000

	Night Charge	Intl Mins	Intl Calls	Intl Charge	CustServ Calls
count	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000
mean	9.039325	10.237294	4.479448	2.764581	1.562856
std	2.275873	2.791840	2.461214	0.753773	1.315491
min	1.040000	0.000000	0.000000	0.000000	0.000000
25%	7.520000	8.500000	3.000000	2.300000	1.000000
50%	9.050000	10.300000	4.000000	2.780000	1.000000
75%	10.590000	12.100000	6.000000	3.270000	2.000000
max	17.770000	20.000000	20.000000	5.400000	9.000000

As was shown phone is all unique values and can be removed from the data frame. Also decided to create dummy variables for the binary attributes.

```
In [6]: df.drop('Phone', axis=1, inplace=True)
```

```
#df.drop('VMail Message', axis=1, inplace=True)
```

```
In [7]: df['Int\l Plan'] = df['Int\l Plan'].map( {'yes': 1, 'no': 0} ).astype(int)
df['VMail Plan'] = df['VMail Plan'].map( {'yes': 1, 'no': 0} ).astype(int)
df['Churn?'] = df['Churn?'].map( {'True.': 1, 'False.': 0} ).astype(int)
df.head()
```

```
Out [7]:
```

	State	Account Length	Area Code	Int'l Plan	VMail Plan	VMail Message \
0	KS	128	415	0	1	25
1	OH	107	415	0	1	26
2	NJ	137	415	0	0	0
3	OH	84	408	1	0	0
4	OK	75	415	1	0	0

	Day Mins	Day Calls	Day Charge	Eve Mins	Eve Calls	Eve Charge	\
0	265.1	110	45.07	197.4	99	16.78	
1	161.6	123	27.47	195.5	103	16.62	
2	243.4	114	41.38	121.2	110	10.30	
3	299.4	71	50.90	61.9	88	5.26	
4	166.7	113	28.34	148.3	122	12.61	

	Night Mins	Night Calls	Night Charge	Intl Mins	Intl Calls	Intl Charge	\
0	244.7	91	11.01	10.0	3	2.70	
1	254.4	103	11.45	13.7	3	3.70	
2	162.6	104	7.32	12.2	5	3.29	
3	196.9	89	8.86	6.6	7	1.78	
4	186.9	121	8.41	10.1	3	2.73	

	CustServ Calls	Churn?
0	1	0
1	1	0
2	0	0
3	2	0
4	3	0


4 D) Are there any outlier values (records) for each of the attributes ?

The box plot shows that some attributes (acc len, calls, mins, charges) have a quite a few outliers. This is most likely due to this data not being a full dataset and the nature of this data is most likely influenced by the pareto principle.

In this case it makes sense to consolidate the fields calls, mins, charges. Also makes sense to bin the data and then see if there is any relation to the class variable.

```
In [8]: #Use boxplot to see if there are any outlier values
df.boxplot(figsize=(20,10))
```

```
Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x24a57f338d0>
```



4.1 E) Analyze the distribution of numeric attributes (normal or other)

As noted above it made sense to consolidate min, calls, and charges. This was done before plotting the distributions

```
In [9]: #Group the day, eve, night data into total calls, total charges and total mins to simp
df['Total Mins Non-Intl'] = df['Day Mins'] + df['Eve Mins'] + df['Night Mins']
df['Total Calls Non-Intl'] = df['Day Calls'] + df['Eve Calls'] + df['Night Calls']
df['Total Charge Non-Intl'] = df['Day Charge'] + df['Eve Charge'] + df['Night Charge']


#Create a new columns for combined totals (non-intl and intl)
df['Total Mins'] = df['Total Mins Non-Intl'] + df['Intl Mins']
df['Total Calls'] = df['Total Calls Non-Intl'] + df['Intl Calls']
df['Total Charge'] = df['Total Charge Non-Intl'] + df['Intl Charge']
#df.head()
```

Account Length - Normally distributed

```
In [10]: #Account Length, Day Mins, Day Charge, Eve Mins, Eve Charge, Night Mins, Night Charge
sns.distplot(df["Account Length"])
```

C:\Users\benpo\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a n
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x24a5deacda0>
```

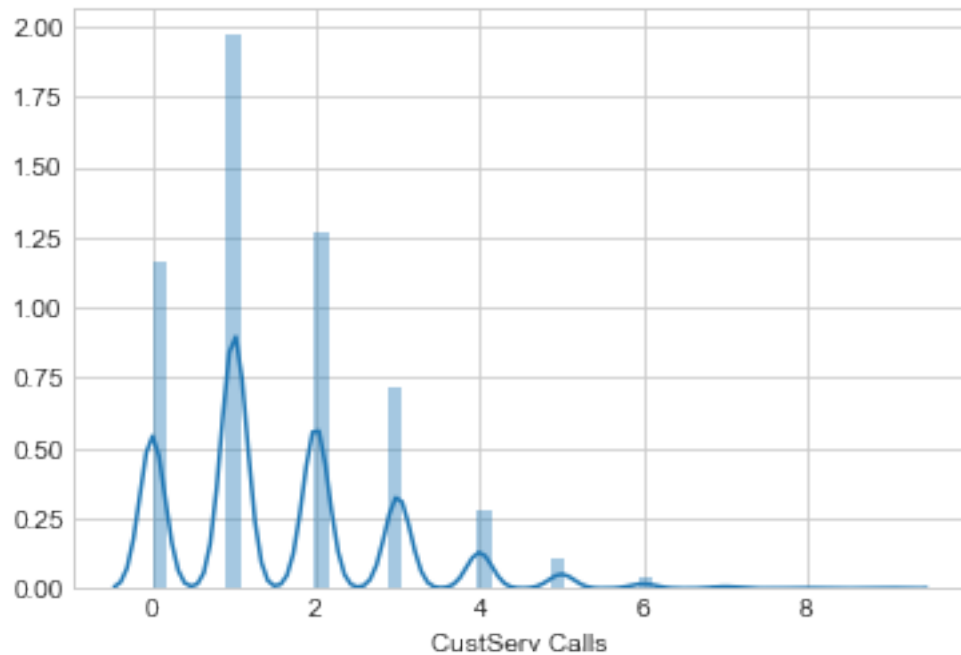


CustServ Calls - Not normally dist and should be converted to a binary value

```
In [11]: sns.distplot(df["CustServ Calls"])
```

```
C:\Users\benpo\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. This will raise an error in a future version.
    return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x24a5db27748>
```



VMail Message - Was not normally dist and should be converted to a binary value. Basically the same as VMail Plan so it was dropped.


```
In [12]: #sns.distplot(df["VMail Message"])
```

Mins - Normally distributed

```
In [13]: sns.distplot(df["Total Mins"])
```

```
C:\Users\benpo\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x24a5dc16b38>
```

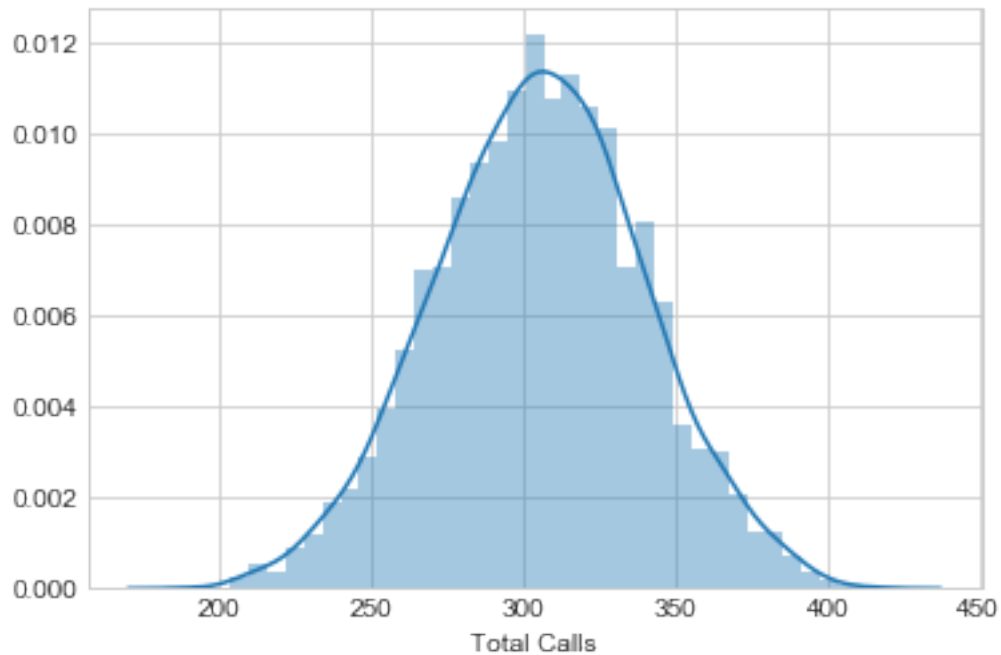



Calls - Normally distributed

```
In [14]: sns.distplot(df["Total Calls"])
```

```
C:\Users\benpo\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will result in a ValueError. Use the tuple form of indexing on new arrays to silence this warning.
    return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x24a5dcd3828>
```




Charge - Normally distributed

```
In [15]: sns.distplot(df["Total Charge"])
```

```
C:\Users\benpo\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will result in an error.
  return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x24a5dd7f978>
```



4.2 F) Analyze attributes and determine if they have any influence on the class

Before starting this analysis the account length was binned. Skipped binning each individual min, call, charge column and focused on the the overall totals.

```
In [16]: #Bin account length
binwidth_al = int((max(df['Account Length'])-min(df['Account Length']))/5)
bins_al = range(min(df['Account Length']), max(df['Account Length']), binwidth_al)
al_names = ['Newest', 'Avg', "Oldest"]
df['Acc Length Binned'] = pd.cut(df['Account Length'], bins_al)

#Convert total mins to int so it can be binned
df['Total Mins'] = df['Total Mins'].astype(np.int64)
#Bin Total Mins
binwidth_tm = int((max(df['Total Mins'])-min(df['Total Mins']))/6)
bins_tm = range(min(df['Total Mins']), max(df['Total Mins']), binwidth_tm)

df['Total Mins Binned'] = pd.cut(df['Total Mins'], bins_tm)

#Convert total Calls to int so it can be binned
df['Total Calls'] = df['Total Calls'].astype(np.int64)
#Bin Total Calls
binwidth_tc = int((max(df['Total Calls'])-min(df['Total Calls']))/6)
bins_tc = range(min(df['Total Calls']), max(df['Total Calls']), binwidth_tc)
df['Total Calls Binned'] = pd.cut(df['Total Calls'], bins_tc)
```

```

#Convert total Calls to int so it can be binned
df['Total Charge'] = df['Total Charge'].astype(np.int64)
#Bin Total Calls
binwidth_tch = int((max(df['Total Charge'])-min(df['Total Charge']))/6)
bins_tch = range(min(df['Total Charge']), max(df['Total Charge']), binwidth_tch)
df['Total Charge Binned'] = pd.cut(df['Total Charge'], bins_tch)

df.head()

```

```

Out[16]:

```

	State	Account Length	Area Code	Int'l Plan	VMail Plan	VMail Message	\
0	KS	128	415	0	1	25	
1	OH	107	415	0	1	26	
2	NJ	137	415	0	0	0	
3	OH	84	408	1	0	0	
4	OK	75	415	1	0	0	

	Day Mins	Day Calls	Day Charge	Eve Mins	...	\
0	265.1	110	45.07	197.4	...	
1	161.6	123	27.47	195.5	...	
2	243.4	114	41.38	121.2	...	
3	299.4	71	50.90	61.9	...	
4	166.7	113	28.34	148.3	...	

	Total Mins Non-Intl	Total Calls Non-Intl	Total Charge Non-Intl	\
0	707.2	300	72.86	
1	611.5	329	55.54	
2	527.2	328	59.00	
3	558.2	248	65.02	
4	501.9	356	49.36	

	Total Mins	Total Calls	Total Charge	Acc Length Binned	\
0	717	303	75	(97, 145]	
1	625	332	59	(97, 145]	
2	539	333	62	(97, 145]	
3	564	255	66	(49, 97]	
4	512	359	52	(49, 97]	

	Total Mins Binned	Total Calls Binned	Total Charge Binned
0	(684, 784]	(302, 339]	(70, 82]
1	(584, 684]	(302, 339]	(58, 70]
2	(484, 584]	(302, 339]	(58, 70]
3	(484, 584]	(228, 265]	(58, 70]
4	(484, 584]	(339, 376]	(46, 58]

[5 rows x 30 columns]

State - Some states are more likely to result in churn but there is no strong correlation visible . It can be dropped.

```
In [17]: df[['State', 'Churn?']].groupby(['State'], as_index=False).mean().sort_values(by='Churn?')
```

```
Out[17]:
```

	State	Churn?
31	NJ	0.264706
4	CA	0.264706
43	TX	0.250000
20	MD	0.242857
40	SC	0.233333
22	MI	0.219178
25	MS	0.215385
33	NV	0.212121
47	WA	0.212121
21	ME	0.209677
26	MT	0.205882
2	AR	0.200000
16	KS	0.185714
34	NY	0.180723
23	MN	0.178571
38	PA	0.177778
19	MA	0.169231
6	CT	0.162162
27	NC	0.161765
30	NH	0.160714
10	GA	0.148148
8	DE	0.147541
36	OK	0.147541
37	OR	0.141026
44	UT	0.138889
5	CO	0.136364
17	KY	0.135593
41	SD	0.133333
35	OH	0.128205
9	FL	0.126984
15	IN	0.126761
13	ID	0.123288
50	WY	0.116883
24	MO	0.111111
46	VT	0.109589
1	AL	0.100000
32	NM	0.096774
28	ND	0.096774
49	WV	0.094340
42	TN	0.094340
7	DC	0.092593
39	RI	0.092308
48	WI	0.089744
14	IL	0.086207
29	NE	0.081967


18	LA	0.078431
12	IA	0.068182
45	VA	0.064935
3	AZ	0.062500
0	AK	0.057692
11	HI	0.056604

Acc Length Binned - No real difference shown. It can be dropped.

```
In [18]: sns.barplot(x='Acc Length Binned', y='Churn?', data=df)
```

C:\Users\benpo\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; using `tuple` instead. Errors may arise from now.
 return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x24a5e1c3be0>
```




Area Code - No predictive power. It can be dropped

```
In [19]: #df[['Area Code', 'Churn?']].groupby(['Area Code'], as_index=False).mean().sort_values(
sns.barplot(x='Area Code', y='Churn?', data=df)
```

C:\Users\benpo\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; using `tuple` instead. Errors may arise from now.
 return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval

```
Out[19]: <matplotlib.axes._subplots.AxesSubplot at 0x24a5e19e6a0>
```




Intl Plan - If customer has international plan, they are more likely to churn. Keep attribute.

```
In [20]: sns.barplot(x='Int\'l Plan', y='Churn?', data=df)
```

```
C:\Users\benpo\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. The warnings are disabled if `np.seterr(all='ignore')` is called.
    return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x24a5e2b1550>
```




VMail Plan - No Voicemail plan does increase chance of churn slightly. Keep for now, but it can be removed later on from model.

```
In [21]: sns.barplot(x='VMail Plan', y='Churn?', data=df)
```

```
C:\Users\benpo\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-bracketed call like np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval is deprecated. Please use np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval instead
```

```
Out[21]: <matplotlib.axes._subplots.AxesSubplot at 0x24a5e2fa828>
```





Total Mins Binned - Minutes don't appear to influence churn as the outcome is fairly random. Will drop field based on this and fact there are a high number of outliers previously identified for this attribute.

```
In [22]: sns.barplot(x='Total Mins Binned', y='Churn?', data=df)
```

```
C:\Users\benpo\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a n
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[22]: <matplotlib.axes._subplots.AxesSubplot at 0x24a5e35b358>
```




Total Calls Binned - Total Calls doesn't appear to influence the class attribute. Will drop field based on this and fact there are a high number of outliers previously identified for this attribute.

```
In [23]: sns.barplot(x='Total Calls Binned', y='Churn?', data=df)
```

```
C:\Users\benpo\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[23]: <matplotlib.axes._subplots.AxesSubplot at 0x24a5e3c0898>
```




Total Charge Binned - Customer is more likely to churn if their charge is high. Create a feature if charge >80 (y/n)

```
In [24]: sns.barplot(x='Total Charge Binned', y='Churn?', data=df)
```

C:\Users\benpo\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. This will raise an error in a future version.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[24]: <matplotlib.axes._subplots.AxesSubplot at 0x24a5e4304a8>
```




```
In [25]: #Create new feature
df['Total Charge > 80'] = df['Total Charge'].apply(lambda x: 1 if x > 80 else 0)
```

CustServ Calls - The more calls equals a much greater chance of churn. Greater than 3 looks like a strong indicator. Create new feature based on this.

```
In [26]: sns.barplot(x='CustServ Calls', y='Churn?', data=df)
```

C:\Users\benpo\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a n
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval

```
Out[26]: <matplotlib.axes._subplots.AxesSubplot at 0x24a5f4399e8>
```



```
In [27]: #Create new feature
df['CustServ Calls >3'] = df['CustServ Calls'].apply(lambda x: 1 if x >3 else 0)
#df.info()
```

4.3 G) Final Clean-up

Above there has already been some data clean up such as removing attributes not needed and converting categorical data to numerical values. In this section anything remaining before data is used for modeling is completed

Remove columns that are not needed (identified in part f)

```
In [28]: #df.drop('State', axis=1, inplace=True)
df.drop('Account Length', axis=1, inplace=True)
df.drop('Area Code', axis=1, inplace=True)
df.drop('Day Mins', axis=1, inplace=True)
df.drop('Day Calls', axis=1, inplace=True)
df.drop('Day Charge', axis=1, inplace=True)
df.drop('Eve Mins', axis=1, inplace=True)
df.drop('Eve Calls', axis=1, inplace=True)
df.drop('Eve Charge', axis=1, inplace=True)
df.drop('Night Mins', axis=1, inplace=True)
df.drop('Night Calls', axis=1, inplace=True)
df.drop('Night Charge', axis=1, inplace=True)
df.drop('Intl Mins', axis=1, inplace=True)
df.drop('Intl Calls', axis=1, inplace=True)
```

```

df.drop('Int'l Charge', axis=1, inplace=True)
df.drop('Total Mins Non-Int'l', axis=1, inplace=True)
df.drop('Total Calls Non-Int'l', axis=1, inplace=True)
df.drop('Total Charge Non-Int'l', axis=1, inplace=True)
df.drop('Total Mins', axis=1, inplace=True)
df.drop('Total Calls', axis=1, inplace=True)
#df.drop('Total Charge', axis=1, inplace=True)
df.drop('Acc Length Binned', axis=1, inplace=True)
df.drop('Total Mins Binned', axis=1, inplace=True)
df.drop('Total Calls Binned', axis=1, inplace=True)
df.drop('Total Charge Binned', axis=1, inplace=True)
df.drop('CustServ Calls', axis=1, inplace=True)

```

There are now 7 features to use for modeling after investigating which attributes can influence the class variable and then removing the ones are unrelated.

```
In [29]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 8 columns):
State                3333 non-null object
Int'l Plan          3333 non-null int32
VMail Plan          3333 non-null int32
VMail Message       3333 non-null int64
Churn?              3333 non-null int32
Total Charge        3333 non-null int64
Total Charge > 80   3333 non-null int64
CustServ Calls >3   3333 non-null int64
dtypes: int32(3), int64(4), object(1)
memory usage: 169.3+ KB

```

Finally for each newly created attribute used determine churn, group by and take a count of no churn.churn to see size of each influencer and confirm it should still be included. Would have excluded if attributes were highly weighted to either no churn or churn.

```
In [30]: df_groupby2 = (df.groupby(['Int\'l Plan']))['Churn?'].agg('count')
df_groupby2
```

```

Out[30]: Int'l Plan
0      3010
1       323
Name: Churn?, dtype: int64

```

```
In [31]: df_groupby2 = (df.groupby(['VMail Plan']))['Churn?'].agg('count')
df_groupby2
```

```
Out[31]: VMail Plan
0      2411
1       922
Name: Churn?, dtype: int64
```

```
In [32]: df_groupby2 = (df.groupby(['Total Charge > 80'])['Churn?'].agg('count'))
df_groupby2
```

```
Out[32]: Total Charge > 80
0      3266
1        67
Name: Churn?, dtype: int64
```

```
In [33]: df_groupby2 = (df.groupby(['CustServ Calls >3'])['Churn?'].agg('count'))
df_groupby2
```

```
Out[33]: CustServ Calls >3
0      3066
1       267
Name: Churn?, dtype: int64
```

4.4 H) Three way data split (training 60%, validation 20%, test 20%)

The function below takes the dataframe and does a three way split based on provided percentges. It also randomizes the index so that the output is a random sample.

```
In [34]: def train_validate_test_split(df, train_percent=.6, validate_percent=.2, seed=None):
    np.random.seed(seed)
    perm = np.random.permutation(df.index)
    m = len(df.index)
    train_end = int(train_percent * m)
    validate_end = int(validate_percent * m) + train_end
    train = df.ix[perm[:train_end]]
    validate = df.ix[perm[train_end:validate_end]]
    test = df.ix[perm[validate_end:]]
    return train, validate, test

np.random.seed([3,1415])

train, validate, test = train_validate_test_split(df)
```

C:\Users\benpo\Anaconda3\lib\site-packages\ipykernel_launcher.py:7: DeprecationWarning: .ix is deprecated. Please use .loc for label based indexing or .iloc for positional indexing

See the documentation here:

<http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprecated>
import sys

```
C:\Users\benpo\Anaconda3\lib\site-packages\ipykernel_launcher.py:8: DeprecationWarning:
.ix is deprecated. Please use
.loc for label based indexing or
.iloc for positional indexing
```

See the documentation here:

<http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprecated>

```
C:\Users\benpo\Anaconda3\lib\site-packages\ipykernel_launcher.py:9: DeprecationWarning:
.ix is deprecated. Please use
.loc for label based indexing or
.iloc for positional indexing
```

See the documentation here:

<http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprecated>

```
if __name__ == '__main__':
```

```
In [35]: test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 668 entries, 373 to 2237
Data columns (total 8 columns):
State                668 non-null object
Int'l Plan           668 non-null int32
VMail Plan           668 non-null int32
VMail Message        668 non-null int64
Churn?               668 non-null int32
Total Charge         668 non-null int64
Total Charge > 80    668 non-null int64
CustServ Calls >3    668 non-null int64
dtypes: int32(3), int64(4), object(1)
memory usage: 39.1+ KB
```

```
In [36]: df.to_csv("customer_churn_processed.csv", sep=',', encoding='utf-8')
```

5 PART 2 - Predictive Modeling (Classification)

6 DECISION TREE

6.1 A) Read partially processed file from earlier analysis and review included variables

```
In [37]: # import dataset
```

```
df = pd.read_csv("customer_churn_processed.csv")
```



```

#Verify import was successful and check for missing values
#Confirmed and there is no missing data
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 9 columns):
Unnamed: 0      3333 non-null int64
State           3333 non-null object
Int'l Plan      3333 non-null int64
VMail Plan      3333 non-null int64
VMail Message   3333 non-null int64
Churn?          3333 non-null int64
Total Charge    3333 non-null int64
Total Charge > 80  3333 non-null int64
CustServ Calls >3  3333 non-null int64
dtypes: int64(8), object(1)
memory usage: 234.4+ KB

```

6.2 B) Confirm relative frequency of churn in the dataset

The results of the analysis show an imbalance of churn versus non-churn

```

In [38]: df_groupby = (df.groupby(['Churn?']))['Churn?'].agg('count')

print (df_groupby)
print ("Churn rate - ",483/(2850+483))

```

```

Churn?
0      2850
1        483
Name: Churn?, dtype: int64
Churn rate -  0.14491449144914492

```

6.3 C) List the names of columns for easy later reference

```

In [39]: # create list with column names

```

```

col_nm = list(df.columns.values)

print (col_nm)

```

```

['Unnamed: 0', 'State', 'Int'l Plan', 'VMail Plan', 'VMail Message', 'Churn?', 'Total Charge',

```

```
In [40]: ## Create dummy variables for state to feed into decision tree
```

```
s = df['State']  
state_dummies = pd.get_dummies(s)
```

6.4 D) Concatenate state_dummies with main df and drop original state variable

```
In [41]: df = pd.concat([df, state_dummies], axis=1, sort=False)  
df.drop('State', axis=1, inplace=True)  
df.drop('Unnamed: 0', axis=1, inplace=True)
```

```
In [42]: # Create train, test and validation data sets (60,20,20)  
# note generates and error but works fine
```

```
def train_validate_test_split(df, train_percent=.6, validate_percent=.2, seed=None):  
    np.random.seed(seed)  
    perm = np.random.permutation(df.index)  
    m = len(df.index)  
    train_end = int(train_percent * m)  
    validate_end = int(validate_percent * m) + train_end  
    train = df.ix[perm[:train_end]]  
    validate = df.ix[perm[train_end:validate_end]]  
    test = df.ix[perm[validate_end:]]  
    return train, validate, test  
  
np.random.seed([3,1415])  
  
train, validate, test = train_validate_test_split(df)
```

```
C:\Users\benpo\Anaconda3\lib\site-packages\ipykernel_launcher.py:10: DeprecationWarning:  
.ix is deprecated. Please use  
.loc for label based indexing or  
.iloc for positional indexing
```

See the documentation here:

<http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprecated>

```
# Remove the CWD from sys.path while we load stuff.
```

```
C:\Users\benpo\Anaconda3\lib\site-packages\ipykernel_launcher.py:11: DeprecationWarning:  
.ix is deprecated. Please use  
.loc for label based indexing or  
.iloc for positional indexing
```

See the documentation here:

<http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprecated>

```
# This is added back by InteractiveShellApp.init_path()
```

```
C:\Users\benpo\Anaconda3\lib\site-packages\ipykernel_launcher.py:12: DeprecationWarning:  
.ix is deprecated. Please use  
.loc for label based indexing or
```

.iloc for positional indexing

See the documentation here:

<http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprecated>
if sys.path[0] == '':

6.5 E) Confirm the split took place correctly.

```
In [43]: # confirm split
print(test.describe())
print(train.describe())
print(validate.describe())
```

	Int'l Plan	VMail Plan	VMail Message	Churn?	Total Charge \
count	668.000000	668.000000	668.000000	668.000000	668.000000
mean	0.109281	0.255988	7.492515	0.149701	59.008982
std	0.312226	0.436742	13.220819	0.357045	10.372951
min	0.000000	0.000000	0.000000	0.000000	25.000000
25%	0.000000	0.000000	0.000000	0.000000	52.000000
50%	0.000000	0.000000	0.000000	0.000000	58.000000
75%	0.000000	1.000000	15.000000	0.000000	66.000000
max	1.000000	1.000000	45.000000	1.000000	92.000000

	Total Charge > 80	CustServ Calls >3	AK	AL \
count	668.000000	668.000000	668.000000	668.000000
mean	0.016467	0.062874	0.014970	0.023952
std	0.127358	0.242919	0.121524	0.153015
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000

	AR	...	SD	TN	TX	UT \
count	668.000000	...	668.000000	668.000000	668.000000	668.000000
mean	0.019461	...	0.023952	0.008982	0.022455	0.025449
std	0.138242	...	0.153015	0.094418	0.148269	0.157603
min	0.000000	...	0.000000	0.000000	0.000000	0.000000
25%	0.000000	...	0.000000	0.000000	0.000000	0.000000
50%	0.000000	...	0.000000	0.000000	0.000000	0.000000
75%	0.000000	...	0.000000	0.000000	0.000000	0.000000
max	1.000000	...	1.000000	1.000000	1.000000	1.000000

	VA	VT	WA	WI	WV	WY
count	668.000000	668.000000	668.000000	668.000000	668.000000	668.000000
mean	0.017964	0.017964	0.019461	0.026946	0.031437	0.025449
std	0.132920	0.132920	0.138242	0.162047	0.174627	0.157603

min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

[8 rows x 58 columns]

	Int'l Plan	VMail Plan	VMail Message	Churn?	Total Charge \
count	1999.000000	1999.000000	1999.000000	1999.000000	1999.000000
mean	0.100050	0.273137	8.034017	0.147074	58.882441
std	0.300142	0.445682	13.724912	0.354268	10.554245
min	0.000000	0.000000	0.000000	0.000000	25.000000
25%	0.000000	0.000000	0.000000	0.000000	52.000000
50%	0.000000	0.000000	0.000000	0.000000	59.000000
75%	0.000000	1.000000	19.000000	0.000000	66.000000
max	1.000000	1.000000	50.000000	1.000000	96.000000

	Total Charge > 80	CustServ Calls >3	AK	AL \
count	1999.000000	1999.000000	1999.000000	1999.000000
mean	0.020510	0.086043	0.014507	0.022511
std	0.141773	0.280498	0.119599	0.148376
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000

	AR	...	SD	TN	TX \
count	1999.000000	...	1999.000000	1999.000000	1999.000000
mean	0.014007	...	0.015508	0.018009	0.024012
std	0.117549	...	0.123592	0.133017	0.153125
min	0.000000	...	0.000000	0.000000	0.000000
25%	0.000000	...	0.000000	0.000000	0.000000
50%	0.000000	...	0.000000	0.000000	0.000000
75%	0.000000	...	0.000000	0.000000	0.000000
max	1.000000	...	1.000000	1.000000	1.000000

	UT	VA	VT	WA	WI \
count	1999.000000	1999.000000	1999.000000	1999.000000	1999.000000
mean	0.020510	0.024512	0.020510	0.019010	0.020510
std	0.141773	0.154672	0.141773	0.136592	0.141773
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

WV

WY

count	1999.000000	1999.000000
mean	0.031516	0.020010
std	0.174751	0.140069
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

[8 rows x 58 columns]

	Int'l Plan	VMail Plan	VMail Message	Churn?	Total Charge \
count	666.000000	666.000000	666.000000	666.000000	666.000000
mean	0.075075	0.307808	8.902402	0.133634	59.147147
std	0.263710	0.461933	14.018391	0.340514	10.489349
min	0.000000	0.000000	0.000000	0.000000	22.000000
25%	0.000000	0.000000	0.000000	0.000000	53.000000
50%	0.000000	0.000000	0.000000	0.000000	59.000000
75%	0.000000	1.000000	22.000000	0.000000	66.000000
max	1.000000	1.000000	51.000000	1.000000	89.000000

	Total Charge > 80	CustServ Calls >3	AK	AL \
count	666.000000	666.000000	666.000000	666.000000
mean	0.022523	0.079580	0.019520	0.028529
std	0.148487	0.270845	0.138446	0.166602
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000

	AR	...	SD	TN	TX	UT \
count	666.000000	...	666.000000	666.000000	666.000000	666.000000
mean	0.021021	...	0.019520	0.016517	0.013514	0.021021
std	0.143562	...	0.138446	0.127547	0.115546	0.143562
min	0.000000	...	0.000000	0.000000	0.000000	0.000000
25%	0.000000	...	0.000000	0.000000	0.000000	0.000000
50%	0.000000	...	0.000000	0.000000	0.000000	0.000000
75%	0.000000	...	0.000000	0.000000	0.000000	0.000000
max	1.000000	...	1.000000	1.000000	1.000000	1.000000

	VA	VT	WA	WI	WV	WY
count	666.000000	666.000000	666.000000	666.000000	666.000000	666.000000
mean	0.024024	0.030030	0.022523	0.028529	0.033033	0.030030
std	0.153239	0.170798	0.148487	0.166602	0.178857	0.170798
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000

```
max      1.000000    1.000000    1.000000    1.000000    1.000000    1.000000
```

```
[8 rows x 58 columns]
```

6.6 F) Create separate target vector to feed into the algorithm.

```
In [44]: # create separate arrays for target variable
```

```
train_target = train['Churn?']
test_target = test['Churn?']
validate_target = validate['Churn?']
```

```
# delete target variable from train and test dataframes
```

```
train = train.drop('Churn?',1)
test = test.drop('Churn?',1)
validate = validate.drop('Churn?',1)
```

6.7 G) Import decision tree and fit model

#importing other algorithms is very similar

```
In [45]: from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import tree
from sklearn.model_selection import train_test_split
```

```
tree = DecisionTreeClassifier(criterion = 'gini', splitter='best', max_depth=None, min
                               min_samples_leaf=10, min_weight_fraction_leaf=0.0, ran
                               max_leaf_nodes=None, min_impurity_decrease=0.0, min_imp
                               presort=False)
```

```
tree.fit (train,train_target)
```

```
Out[45]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=10, min_samples_split=5,
                                min_weight_fraction_leaf=0.0, presort=False, random_state=0,
                                splitter='best')
```

```
In [46]: # Score train data set
```

```
In [47]: # Output model scores for test
```

```
test_pred = tree.predict(test)
```

```
In [48]: # Calculate model evaluation metrics
```

```
from sklearn.metrics import precision_recall_fscore_support as score
from sklearn.metrics import accuracy_score
precision, recall, fscore, support = score(test_target, test_pred)
```

```
precision = ('precision: {}'.format(precision))
recall = ('recall: {}'.format(recall))
fscore = ('fscore: {}'.format(fscore))
support = ('support: {}'.format(support))
accuracy = accuracy_score(test_target, test_pred, normalize=True)

output = {accuracy, precision, recall, fscore, support}
output
```

```
Out[48]: {0.938622754491018,
          'fscore: [0.96480687 0.76023392]',
          'precision: [0.94137353 0.91549296]',
          'recall: [0.98943662 0.65      ]',
          'support: [568 100]}'
```

```
In [49]: # Create confusion matrix
```

```
In [50]: df_confusion = pd.crosstab(test_target, test_pred)
print (df_confusion)
```

```
col_0    0    1
Churn?
0        562    6
1         35   65
```

```
In [51]: # Use grid search to improve performance
```

```
In [52]: from sklearn.model_selection import GridSearchCV
#from sklearn.model_selection import GridSearchCV
```

```
parameters = {'min_samples_split':np.arange(2, 80), 'max_depth': np.arange(2,10), 'cr
tree = DecisionTreeClassifier()
grid = GridSearchCV(tree, parameters,scoring='accuracy', cv=8)

grid.fit(train, train_target)
print('The parameters combination that would give best accuracy is : ')
print(grid.best_params_)
print('The best accuracy achieved after parameter tuning via grid search is : ', grid
```

The parameters combination that would give best accuracy is :
{'criterion': 'gini', 'max_depth': 3, 'min_samples_split': 37}
The best accuracy achieved after parameter tuning via grid search is : 0.9464732366183092

```
In [53]: # enter optimized hyperparameters
```

```
In [54]: tree2 = DecisionTreeClassifier(criterion = 'gini', splitter='best', max_depth=3, min_
min_samples_leaf=5, min_weight_fraction_leaf=0.0, rand
max_leaf_nodes=None, min_impurity_decrease=0.0, min_imp
presort=False)
```

```
tree2.fit (train,train_target)
```

```
Out [54]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=3,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=5, min_samples_split=5,
min_weight_fraction_leaf=0.0, presort=False, random_state=0,
splitter='best')
```

```
In [55]: #rescore test with updated hyperparameters
```

```
test_pred2 = tree2.predict(test)
```

```
In [56]: #recheck metrics based on grid search based hyperparameter settings
```

```
precision, recall, fscore, support = score(test_target, test_pred2)
```

```
precision = ('precision: {}'.format(precision))
recall = ('recall: {}'.format(recall))
fscore = ('fscore: {}'.format(fscore))
support = ('support: {}'.format(support))
accuracy = accuracy_score(test_target, test_pred2, normalize=True)
```

```
output = {accuracy, precision, recall, fscore, support}
output
```

```
Out [56]: {0.9446107784431138,
'fscore: [0.96834902 0.77844311]',
'precision: [0.94176373 0.97014925]',
'recall: [0.99647887 0.65      ]',
'support: [568 100]}'
```

```
In [57]: # create confusion matrix for optimized algorithm
```

```
df_confusion = pd.crosstab(test_target, test_pred2)
print (df_confusion)
```



```
col_0      0    1
Churn?
0          566   2
1           35  65
```

In [58]: *# code below to check model performance on validate - not currently working*

```
In [59]: estimator = tree2.fit (train,train_target)
        validate_pred2 = tree2.predict(validate)
        precision, recall, fscore, support = score(validate_target, validate_pred2)

        precision = ('precision: {}'.format(precision))
        recall = ('recall: {}'.format(recall))
        fscore = ('fscore: {}'.format(fscore))
        support = ('support: {}'.format(support))
        accuracy = accuracy_score(test_target, test_pred2, normalize=True)

        final_output = {accuracy, precision, recall, fscore, support}
        final_output
```

```
Out[59]: {0.9446107784431138,
          'fscore: [0.97792869 0.83116883]',
          'precision: [0.95840266 0.98461538]',
          'recall: [0.9982669 0.71910112]',
          'support: [577 89]}'
```

7 NAIVE BAYES

Now on to the second classification model naive bayes

```
In [60]: #verify categorical data converted to numeric
        df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 58 columns):
Int'l Plan      3333 non-null int64
VMail Plan      3333 non-null int64
VMail Message   3333 non-null int64
Churn?          3333 non-null int64
Total Charge    3333 non-null int64
Total Charge > 80 3333 non-null int64
CustServ Calls >3 3333 non-null int64
AK              3333 non-null uint8
AL              3333 non-null uint8
AR              3333 non-null uint8
AZ              3333 non-null uint8
```

CA	3333 non-null uint8
CO	3333 non-null uint8
CT	3333 non-null uint8
DC	3333 non-null uint8
DE	3333 non-null uint8
FL	3333 non-null uint8
GA	3333 non-null uint8
HI	3333 non-null uint8
IA	3333 non-null uint8
ID	3333 non-null uint8
IL	3333 non-null uint8
IN	3333 non-null uint8
KS	3333 non-null uint8
KY	3333 non-null uint8
LA	3333 non-null uint8
MA	3333 non-null uint8
MD	3333 non-null uint8
ME	3333 non-null uint8
MI	3333 non-null uint8
MN	3333 non-null uint8
MO	3333 non-null uint8
MS	3333 non-null uint8
MT	3333 non-null uint8
NC	3333 non-null uint8
ND	3333 non-null uint8
NE	3333 non-null uint8
NH	3333 non-null uint8
NJ	3333 non-null uint8
NM	3333 non-null uint8
NV	3333 non-null uint8
NY	3333 non-null uint8
OH	3333 non-null uint8
OK	3333 non-null uint8
OR	3333 non-null uint8
PA	3333 non-null uint8
RI	3333 non-null uint8
SC	3333 non-null uint8
SD	3333 non-null uint8
TN	3333 non-null uint8
TX	3333 non-null uint8
UT	3333 non-null uint8
VA	3333 non-null uint8
VT	3333 non-null uint8
WA	3333 non-null uint8
WI	3333 non-null uint8
WV	3333 non-null uint8
WY	3333 non-null uint8

dtypes: int64(7), uint8(51)

memory usage: 348.4 KB

7.1 A) Gaussian Model

```
In [61]: #import library for Gaussian Naive Bayes
        from sklearn.naive_bayes import GaussianNB, BernoulliNB

        #set the classifier
        gnb = GaussianNB()

        #set features
        used_features = ["Int'l Plan", "VMail Plan", "Total Charge > 80", "CustServ Calls >3"]

In [62]: #train classifier
        gnb.fit(train[used_features].values, train_target)

        response_var = gnb.predict(validate[used_features])

        #print results

        print("Number of mislabeled points out of a total {} points : {}, performance {:.05.2f}f".format(
            test.shape[0],
            (validate_target != response_var).sum(),
            100*(1-(validate_target != response_var).sum()/test.shape[0])))
```

Number of mislabeled points out of a total 668 points : 91, performance 86.38%.

```
In [63]: #import libraries for model evaluation
        from sklearn.metrics import classification_report
        from sklearn.metrics import confusion_matrix
        from sklearn.metrics import accuracy_score
        from sklearn import metrics

In [64]: #model evaluation
        print(metrics.classification_report(validate_target, response_var))
```

	precision	recall	f1-score	support
0	0.94	0.90	0.92	577
1	0.49	0.61	0.54	89
micro avg	0.86	0.86	0.86	666
macro avg	0.71	0.75	0.73	666
weighted avg	0.88	0.86	0.87	666

```
In [65]: #confusion matrix
print(metrics.confusion_matrix(validate_target, response_var))
```

```
[[521  56]
 [ 35  54]]
```


```
In [66]: # consider the distribution of the features
```

```
sns.distplot(df["Int'l Plan"])
```

```
C:\Users\benpo\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; using `tuple` instead. Errors may arise from you using this form of indexing, so please switch to proper indexing. See https://numpy.org/doc/stable/release/1.13.0-bugfixes.html for more.
```

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[66]: <matplotlib.axes._subplots.AxesSubplot at 0x24a5f66afd0>
```




```
In [67]: sns.distplot(df["VMail Plan"])
```

```
C:\Users\benpo\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; using `tuple` instead. Errors may arise from you using this form of indexing, so please switch to proper indexing. See https://numpy.org/doc/stable/release/1.13.0-bugfixes.html for more.
```

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```


```
Out[67]: <matplotlib.axes._subplots.AxesSubplot at 0x24a5f640438>
```



```
In [68]: sns.distplot(df["Total Charge > 80"])
```

```
C:\Users\benpo\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```


```
Out[68]: <matplotlib.axes._subplots.AxesSubplot at 0x24a5f4f8f60>
```



```
In [69]: sns.distplot(df["CustServ Calls >3"])
```

```
C:\Users\benpo\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. This will raise an error in future versions.
  return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[69]: <matplotlib.axes._subplots.AxesSubplot at 0x24a5f6f0630>
```



7.2 B) Bernoulli

Our selected features are 1s and 0s. Lets try the Bernoulli model.

```
In [70]: #set the classifier
        bnb = BernoulliNB()

        bnb.fit(train[used_features].values, train_target)

        bnb_response_var = bnb.predict(validate[used_features])

        #print results
        print("Number of mislabeled points out of a total {} points : {}, performance {:.05.2f}
              .format(
                  test.shape[0],
```

```
(validate_target != bnb_response_var).sum(),
100*(1-(validate_target != bnb_response_var).sum()/test.shape[0])))
```

Number of mislabeled points out of a total 668 points : 80, performance 88.02%.

In [71]: *#model evaluation*

```
print(metrics.classification_report(validate_target, bnb_response_var))
```

	precision	recall	f1-score	support
0	0.88	0.99	0.93	577
1	0.74	0.16	0.26	89
micro avg	0.88	0.88	0.88	666
macro avg	0.81	0.57	0.60	666
weighted avg	0.86	0.88	0.84	666

In [72]: *#confusion matrix*

```
print(metrics.confusion_matrix(validate_target, bnb_response_var))
```

```
[[572  5]
 [ 75 14]]
```


7.3 C) Model Selection

In [74]: *#Gaussian Classification Report*

```
from yellowbrick.classifier import ClassificationReport
```

```
visualizer = ClassificationReport(gnb,support=True)
```


```
visualizer.fit(train[used_features], train_target) # Fit the visualizer and the model
visualizer.score(validate[used_features], validate_target) # Evaluate the model on the test set
g = visualizer.poof()
```



```
In [75]: #Bernoulli Classification Report
from yellowbrick.classifier import ClassificationReport

visualizer = ClassificationReport(bnb,support=True)

visualizer.fit(train[used_features], train_target) # Fit the visualizer and the model
visualizer.score(validate[used_features], validate_target) # Evaluate the model on the validation data
g = visualizer.poof()
```





7.4 D) Bernoulli and the Test Set

```
In [76]: #Classification Report
from yellowbrick.classifier import ClassificationReport

visualizer = ClassificationReport(bnb,support=True)

visualizer.fit(train[used_features], train_target) # Fit the visualizer and the model
visualizer.score(test[used_features], test_target) # Evaluate the model on the test set
g = visualizer.poof()
```



```
In [77]: bnb.fit(train[used_features].values, train_target)

        bnb_predicted_test = bnb.predict(test[used_features])

        #print results
        print("Number of mislabeled points out of a total {} points : {}, performance {:.05.2f}%".format(
            test.shape[0],
            (test_target != bnb_predicted_test).sum(),
            100*(1-(test_target != bnb_predicted_test).sum()/test.shape[0])))
```

Number of mislabeled points out of a total 668 points : 93, performance 86.08%.

8 RANDOM FOREST

The third and final model considered was random forest.

```
In [78]: from sklearn.ensemble import RandomForestClassifier
```

8.1 A) Initilize Random Forest Classifier & Initial fit

```
In [79]: rf = RandomForestClassifier(n_estimators=100, criterion='gini', max_depth=3, min_samples_leaf=50, min_weight_fraction_leaf=0.0, max_features='auto', max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=None, random_state=42, verbose=0, warm_start=False, class_weight=None)
```

```
In [80]: rf.fit(train, train_target)
```

```
Out[80]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini', max_depth=3, max_features='auto', max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=50, min_samples_split=19, min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=None, oob_score=False, random_state=42, verbose=0, warm_start=False)
```

```
In [81]: test_pred = rf.predict(test)
```

8.2 B) Instantiate performace metrics

```
In [82]: precision, recall, fscore, support = score(test_target, test_pred)
precision = ('precision: {}'.format(precision)),
recall = ('recall: {}'.format(recall)),
fscore = ('fscore: {}'.format(fscore)),
support = ('support: {}'.format(support)),
accuracy = accuracy_score(test_target, test_pred, normalize=True)
output = {accuracy, precision, recall, fscore, support}
output
```

```
C:\Users\benpo\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1143: UndefinedMetricWarning: Precision is not defined because no predicted samples belong to the specified class: 'precision', 'predicted', average, warn_for)
```

```
Out[82]: {('fscore: [0.91909385 0.          ]',),
('precision: [0.8502994 0.          ]',),
('recall: [1. 0.]',),
('support: [568 100]',),
0.8502994011976048}
```

```
In [83]: df_confusion = pd.crosstab(test_target, test_pred)
print (df_confusion)
```

```
col_0    0
Churn?
0        568
1        100
```

```
In [84]: from sklearn.model_selection import GridSearchCV
```

```
In [85]: precision, recall, fscore, support = score(test_target, test_pred)
precision = ('precision: {}'.format(precision))
recall = ('recall: {}'.format(recall))
fscore = ('fscore: {}'.format(fscore))
support = ('support: {}'.format(support))
accuracy = accuracy_score(test_target, test_pred, normalize=True)

output = {accuracy, precision, recall, fscore, support}
output
```

C:\Users\benpo\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1143: UndefinedMetricWarning: Precision is not defined because no predicted samples were equal to the true labels.

```
'precision', 'predicted', average, warn_for)
```

```
Out[85]: {0.8502994011976048,
          'fscore: [0.91909385 0.          ]',
          'precision: [0.8502994 0.          ]',
          'recall: [1. 0.]',
          'support: [568 100]'}

```

```
In [86]: print(df_confusion)
```

```
col_0      0
Churn?
0         568
1         100
```

8.3 C) Grid Search for optimized parameters

```
In [87]: parameters = {'n_estimators': [500,1000,1500,2000],
                        'max_features': ['auto', 'sqrt', 'log2', .9, .2],
                        'max_depth': [2,4,6,8,10],
                        'criterion':['gini','entropy'],
                        # 'bootstrap': [True, False],
                        'min_samples_leaf': [1,2,3]
                        }

rand_for = RandomForestClassifier()
grid = GridSearchCV(rf, parameters, scoring='accuracy', cv=4)
grid.fit(train, train_target)
```

```
Out[87]: GridSearchCV(cv=4, error_score='raise-deprecating',
                      estimator=RandomForestClassifier(bootstrap=True, class_weight=None, criterion=
                      max_depth=3, max_features='auto', max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=50, min_samples_split=19,
                      min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=None,
                      oob_score=False, random_state=42, verbose=0, warm_start=False),
                      fit_params=None, iid='warn', n_jobs=None,
```

```
param_grid={'max_features': ['auto', 'sqrt', 'log2', 0.9, 0.2], 'max_depth': [
pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
scoring='accuracy', verbose=0)
```

The optimized parameters are:

```
In [88]: print(grid.best_params_)
         print(grid.best_score_)
```

```
{'criterion': 'gini', 'max_depth': 4, 'max_features': 0.9, 'min_samples_leaf': 1}
0.9449724862431216
```

8.4 D) Insert new parameters. Fit & Train Model

```
In [102]: rf = RandomForestClassifier(n_estimators=2000, criterion='gini', max_depth=6, min_sam
min_samples_leaf=2, min_weight_fraction_leaf=0.0, max_features=0.9, max_leaf_n
min_impurity_decrease=0.0, min_impurity_split=None, bootstrap='auto', oob_sco
n_jobs=None, random_state=42, verbose=0, warm_start=False, class_weight=None)
rf.fit(train,train_target)
```

```
Out[102]: RandomForestClassifier(bootstrap='auto', class_weight=None, criterion='gini',
max_depth=6, max_features=0.9, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=2, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=2000, n_jobs=None,
oob_score=False, random_state=42, verbose=0, warm_start=False)
```

```
In [103]: rf_pred = rf.predict(test)
```

8.5 E) Re-test performance metrics

```
In [104]: precision, recall, fscore, support = score(test_target, rf_pred)
```

```
precision = ('precision: {}'.format(precision))
recall = ('recall: {}'.format(recall))
fscore = ('fscore: {}'.format(fscore))
support = ('support: {}'.format(support))
accuracy = accuracy_score(test_target, rf_pred, normalize=True)

output = {accuracy, precision, recall, fscore, support}
output
```

```
Out[104]: {0.9446107784431138,
'fscore: [0.96829477 0.78106509]',
'precision: [0.94323873 0.95652174]',
'recall: [0.99471831 0.66      ]',
'support: [568 100]}'
```

```
In [105]: df_confusion = pd.crosstab(test_target, rf_pred)
          print (df_confusion)
```

```
col_0    0    1
Churn?
0         565    3
1          34   66
```

8.6 F) Evaluate performance on validation set

```
In [106]: rf_pred = rf.predict(validate)
          precision, recall, fscore, support = score(validate_target, rf_pred)
          precision = ('precision: {}'.format(precision))
          recall = ('recall: {}'.format(recall))
          fscore = ('fscore: {}'.format(fscore))
          support = ('support: {}'.format(support))
          accuracy = accuracy_score(validate_target, rf_pred, normalize=True)
          output = {accuracy, precision, recall, fscore, support}
          output
```

```
Out[106]: {0.960960960960961,
           'fscore: [0.97792869 0.83116883]',
           'precision: [0.95840266 0.98461538]',
           'recall: [0.9982669 0.71910112]',
           'support: [577 89]}'
```

```
In [107]: df_confusion = pd.crosstab(validate_target, rf_pred)
          print (df_confusion)
```

```
col_0    0    1
Churn?
0         576    1
1          25   64
```

8.7 G) Determine most important features

```
In [108]: importances = rf.feature_importances_
          std = np.std([tree.feature_importances_ for tree in rf.estimators_],
                        axis=0)
          indices = np.argsort(importances)[::-1]

          print("Feature ranking:")

          for f in range (5):
              print("%d. feature %d (%f)" % (f + 1, indices[f], importances[indices[f]]))
```

Feature ranking:

1. feature 3 (0.555415)
2. feature 5 (0.153025)
3. feature 0 (0.105762)
4. feature 2 (0.084555)
5. feature 1 (0.062942)

```
In [109]: train.columns
```

```
Out[109]: Index(['Int'l Plan', 'VMail Plan', 'VMail Message', 'Total Charge',  
                'Total Charge > 80', 'CustServ Calls >3', 'AK', 'AL', 'AR', 'AZ', 'CA',  
                'CO', 'CT', 'DC', 'DE', 'FL', 'GA', 'HI', 'IA', 'ID', 'IL', 'IN', 'KS',  
                'KY', 'LA', 'MA', 'MD', 'ME', 'MI', 'MN', 'MO', 'MS', 'MT', 'NC', 'ND',  
                'NE', 'NH', 'NJ', 'NM', 'NV', 'NY', 'OH', 'OK', 'OR', 'PA', 'RI', 'SC',  
                'SD', 'TN', 'TX', 'UT', 'VA', 'VT', 'WA', 'WI', 'WV', 'WY'],  
                dtype='object')
```

We see that the following feature ranking of importance for the Random Forest: 1. Total Charge (55.54%) 2. CustServ Calls >3 (15.30%) 3. Int'l Plan (10.58%) 4. VMail Messages (8.45%) 5. VMail Plan (6.29)